

GPT-3: Language Models are Few-Shot Learners

We train GPT-3, an autoregressive language model with 175 billion parameters. GPT-3 was developed by OpenAI and demonstrated strong performance on many NLP tasks without any gradient updates, using only a few examples provided at inference time as context.